Inpatient claim review assistant

The system in one sentence

Help claims staff check inpatient cover and charges, with proof for each recommendation.

Key design calls

Staff decide every claim. Rejected: Automatic settlement. Why: Rules need insurer approval.

Code sets eligible charges. Rejected: Model maths. Why: Bills need exact checks.

Use the right policy edition. Rejected: Newest upload. Why: Cover changes over time.

Check the saved review. Rejected: Blind retry. Why: Avoid duplicate records.

Who acts and who decides

Action	Autonomy
Read and cite documents	suggests
Check policy charges	acts alone on approved rules
Save and assign a draft	acts alone after access checks
Approve, reject or pay	human decides

Top risks by RPN

FM3 | Room limit and maths error | RPN 225

FM2 | Wrong policy edition | RPN 180

FM4 | Injected or contradictory evidence | RPN 160

FM6 | Unknown save or false completion | RPN 160

FM7 | Review queue stalls | RPN 160

Launch blockers and owners

H1, H5: First release creates review packets...; Review queue acknowledges assignment within... Owner: Claims lead.

H2: Insurer supplies immutable policy editions... Owner: Policy owner.

H3: Adapters support tenant-scoped reads and... Owner: Platform owner.

H4: Medical documents stay in an... Owner: Privacy owner.

Illustrative cost per case

USD 0.03780 per case in the model budget. This is an estimate, not a current provider quote. Staff and tools cost extra.

Review status: self-review; round 2; pass.

Proposed design. Full rules and blockers: design.json and assumptions.md.

Task tree

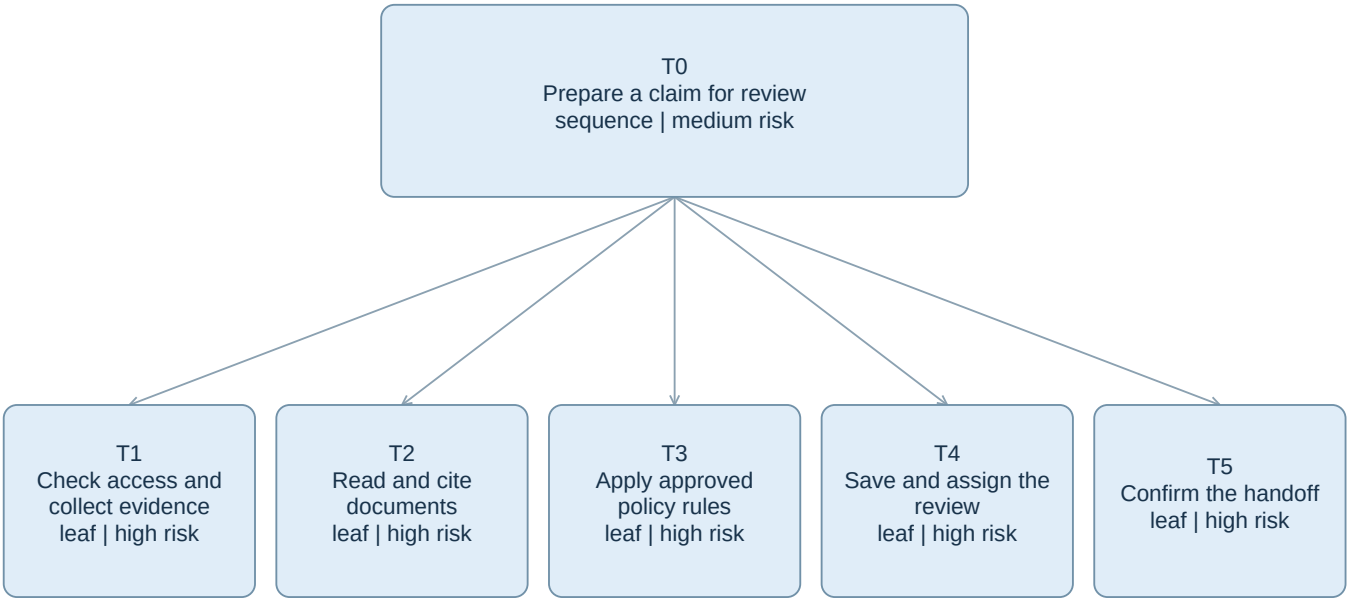

Labels show each leaf mode and risk. Full task inputs, examples and limits are in design.json and PRD.md.

Plain words and scope

First release prepares and assigns reviews. Staff resolve all claims. Missing facts stay pending. No medical advice is given.

A fingerprint is a check that the approved details have not changed. A record ID is proof from a tool that the change was saved.

Workflow diagram

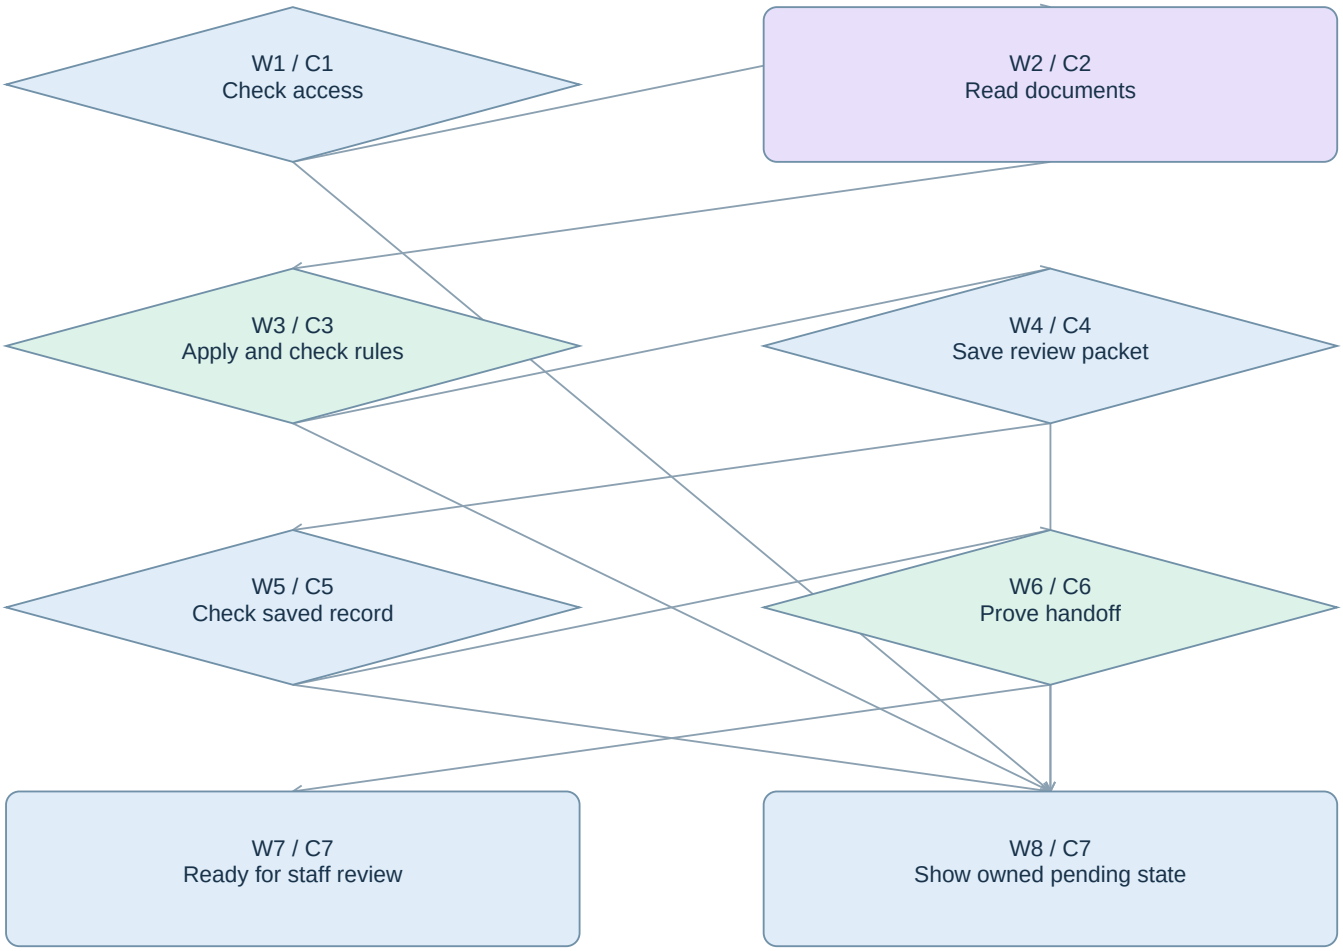

Diamonds mark choices. Arrows show all routes, including retry and recovery. Exact branch rules are in architecture.md.

Colour	Role
#E8E0FA	Agents (LLM)
#E0EDF8	Code: tools
#DDF2E9	Code: validation
#FFF0D5	Code: human triggers

Top risks and release checks

Risk heat table - severity x occurrence

Each cell names the risks at that severity and occurrence. The full scores, causes, controls and owners are in fmea-register.csv.

S / O	1	2	3	4	5	6	7	8	9	10
10	-	-	FM1	-	-	-	-	-	-	-
9	-	-	-	FM2	FM3	-	-	-	-	-
8	-	-	-	FM5, FM6, FM7	FM4	-	-	-	-	-
7	-	-	-	-	-	-	-	-	-	-
6	-	-	-	-	-	-	-	-	-	-
5	-	-	-	-	-	-	-	-	-	-
4	-	-	-	-	-	-	-	-	-	-
3	-	-	-	-	-	-	-	-	-	-
2	-	-	-	-	-	-	-	-	-	-
1	-	-	-	-	-	-	-	-	-	-

Risk	Control summary / full register linked
FM1 Wrong person or privacy leak	Authenticate before fetch; subject-match pages; regional redaction; no text logs.
FM2 Wrong policy edition	Versioned registry date selection; code rechecks versions; missing or conflicting edition means pending.
FM3 Room limit and maths error	Explicit charge-class rules and integer minor-unit arithmetic; missing rules means review.
FM4 Injected or contradictory evidence	Treat all pages as data; citations/schema; model facts never trusted; ambiguity routes staff.
FM5 Duplicate or stale save	Server conditional versions and unique operation key; no final financial effects.
FM6 Unknown save or false completion	Zero write retry; bounded readback, body comparison, record proof, durable on-call recovery.
FM7 Review queue stalls	Atomic assigned reviewer required; queue SLA and business-calendar alerts; never mark claim approved.

Risk scores are proposed estimates. They are not measured odds or risk acceptance.

Components and contracts

This inventory is a summary. Exact typed inputs, outputs and permission rules are in architecture.md and the build milestone files.

Component / group	Responsibility	Task IDs
C1 / Code: tools	Trusted case reader	T1
C2 / Agents (LLM)	Evidence reader	T2
C3 / Code: validation	Rule and packet validator	T3, T4
C4 / Code: tools	Review packet store	T4
C5 / Code: tools	Readback and reconciliation	T4, T5
C6 / Code: validation	Record proof check	T5
C7 / Code: tools	Read-only status view	T5

Shared control boundary

Code checks identity, money, dates and limits. A model may suggest a choice. It cannot authorise or perform a business write.

Architecture and model budget

F1 - Code owns the workflow

Durable code state machine plus one bounded document reader A single evidence reader fits judgement; persistence, rule maths and recovery remain code.

Why: Keeps evidence, authority and recovery explicit.

F2 - Keep long-lived state durable

Database-backed worker with transactional outbox and versioned run state Compare state machine and durable workflow engine on recovery, version migrations, audit and operations; choose code worker for one bounded pipeline. No current vendor capability is asserted.

Why: Keeps evidence, authority and recovery explicit.

F3 - Limit service permissions

Reader can read one authorized claim; writer can create one draft packet; recovery can only read operation status Fresh role/version check at server write; model has no credentials; assignment is provisional and does not approve claim.

Why: Keeps evidence, authority and recovery explicit.

F4 - Keep medical records separate

Regional encrypted document store plus relational authoritative policy/claim ledger; retrieval index contains redacted citations only Privacy owner approves access/region/retention. Encrypt with scoped keys; audit record IDs and access actors without medical text. Delete model cache immediately; proposed retention...

Why: Keeps evidence, authority and recovery explicit.

Model cost

Illustrative model cost is USD 0.03780 per case. Unknown prices are not free service. Confirm rates, traffic, staff and tools before launch.

Full call counts, token budgets, rejected choices and costs: design.json, architecture.md and assumptions.md.

Evaluation coverage matrix

Risk / task	EV1	EV2	EV3	EV4	EV5	EV6	EV7	EV8	EV9	EV10	EV11	EV12
FM1	-	X	-	-	-	-	-	-	-	-	X	-
FM2	-	-	X	-	-	-	-	-	-	-	-	X
FM3	X	-	-	X	-	-	-	-	-	-	-	-
FM4	-	-	-	-	X	-	-	-	-	-	-	-
FM5	-	-	-	-	-	X	X	-	-	-	-	X
FM6	-	-	-	-	-	-	-	X	X	X	-	-
FM7	-	-	-	-	-	-	-	-	-	-	-	-
T1	X	X	-	-	-	-	-	-	-	-	X	-
T2	X	-	-	-	X	-	-	-	-	-	X	-
T3	X	-	X	X	X	-	-	-	-	-	-	X
T4	X	-	-	-	-	X	X	X	-	X	-	X
T5	X	-	-	-	-	-	-	X	X	X	-	-

Evaluation coverage matrix - continued

Risk / task	EV13	EV14	EV15	EV16	EV17
FM1	-	-	-	-	-
FM2	-	-	-	X	X
FM3	X	-	-	-	-
FM4	-	X	-	X	X
FM5	-	-	-	-	-
FM6	-	-	-	-	-
FM7	-	-	X	-	-
T1	-	-	-	-	-
T2	-	X	-	X	X
T3	X	X	-	X	X
T4	-	-	X	X	-
T5	-	-	X	-	-

X means the test is linked to that risk or task. It does not mean the test passed. Complete fixtures and assertions: eval-set.json.

Evaluation plan

Test	Layer	Purpose
EV1	golden	Valid room cap draft
EV2	golden	Cross-tenant rejection
EV3	golden	Historical rule selection
EV4	golden	Missing proration rule routes review
EV5	golden	Injected hospital page
EV6	golden	Concurrent duplicate
EV7	golden	Revoked role at commit
EV8	golden	Save committed but response lost
EV9	golden	Missing proof cannot complete
EV10	golden	Unknown remains owned and frozen
EV11	golden	Raw medical text kept out
EV12	golden	Concurrent policy change
EV13	golden	Minor-unit cap boundary
EV14	golden	Unverified facts and model outage
EV15	golden	Unassigned queue and overdue ownership
EV16	golden	Exclusion is a review candidate
EV17	golden	Policy date boundary conflict
EV18	offline	Held-out evidence and rules
EV19	online	Operational safety monitor
EV20	release	All-case release gate

These are test plans. No business controls, provider quality or user outcomes were measured. See eval-set.json for fixtures and pass rules.

Measurement and rollout

R1 - Reduce reviewer effort

Baseline then assist. No improvement claim until powered comparison completes.

Why: Keeps evidence, authority and recovery explicit.

R2 - Protect claimants

All phases. Zero safety incidents; separate strata quality review.

Why: Keeps evidence, authority and recovery explicit.

R3 - Shadow before staff use

Shadow: proposed 2 weeks minimum. Entry: H2/H4 approved test use and deidentified fixtures. Exit: EV1-EV18 run in harness; SME disagreement resolved; adapter spike H3 complete.

Why: Keeps evidence, authority and recovery explicit.

R4 - Assist with draft packets

Assist: proposed 2 weeks minimum. Entry: H1-H5 signed and EV20 passes; trained reviewers. Exit: all assignments owned, no safety incident, representative queue SLA evidence and complete metrics. Stop: disable inserts; retain readback/recovery.

Why: Keeps evidence, authority and recovery explicit.

R5 - Bounded pilot then scale assist

Pilot: randomized assist-only pilot for >=4 weeks after power sizing; scale remains assist-only.. Entry: assist exit, approved sample/power plan and owner authorization. Exit: powered R1 benefit, no safety breach, no worsened audited decision rate and H5 capacity. Scale...

Why: Keeps evidence, authority and recovery explicit.

R6 - Compare reviewer effort fairly

Pilot. Collect baseline variance and intrareviewer correlation, power a two-sided mean comparison at proposed alpha .05/power .8; adjust reviewer-cluster design effect and attrition. Unknown variance prevents sample claim.

Why: Keeps evidence, authority and recovery explicit.

R7 - Own failures and recovery

All phases. Signed runbooks, writer kill switch drill, pending-operation recovery drill and queue coverage before launch; only insurer executives may accept residual risk. No acceptance recorded here.

Why: Keeps evidence, authority and recovery explicit.

Build handoff and review

Start with build/00-overview.md. Load one milestone file at a time. Exact source rules remain in design.json; do not treat design approval as live write permission.

X1: Implement only the bounded inpatient review-packet workflow. Preserve evidence gates and draft-only permissions.

Review dimensions - evidence in review.json

Dimension	Score	Dimension	Score
D1	3/4	D7	3/4
D2	3/4	D8	3/4
D3	3/4	D9	3/4
D4	3/4	D10	3/4
D5	3/4	D11	3/4
D6	3/4	D12	3/4

ISSUE1 major resolved: Initial artifact left descriptions empty on many cards; validator rejects incomplete cards. Fill specific summaries and revalidate. Repaired and inspected in current design; validator...

ISSUE2 major resolved: W8 lists T1/T4 but C7 serves only T5. Align display responsibility with T5. Repaired and inspected in current design; validator now passes.

ISSUE3 major resolved: Missing-proof fixture says start after prior operation but includes a POST in mocks. Explicitly start proof stage and remove POST to isolate guard. Repaired...

Glossary

canonical: agreed

discriminated: typed

payload binding: approved details

digest: fingerprint

idempotency: duplicate-write protection

reciprocal: two-way

terminal disposition: final outcome

calc_digest: a fingerprint of the approved amount

After delivering DESIGN.pdf, ask the user: Would you also like a visual HTML version in the same ten-stage structure as the app? Generate HTML only after opt-in, unless already requested.

Full assumptions, owners and evidence: assumptions.md. Review method and failed rounds are preserved in review.json and review-round-*.json.